

CS 779 Advance Database Management  
Term Project

# Covid19 social impact mining by Spark

Student: Chen Wu

Professor: Jack Polnar

# Description

Covid19 has been heavily impacting social life this year, we could see changes on every social aspect. This project measures public safety, energy usage changes associated to covid19.

- Goal: Explore social changes caused by Covid-19 this year.
- How to: collecting data from New York Government, specific in Covid-19 data, crime, utility usage, vehicle collisions and shooting cases.
- Measure: To see how those data change by month in 2020.
- Tool: Spark on Databricks
- Language: SQL, Python

## Data collection

I collected data from New York City public data: <https://opendata.cityofnewyork.us/>

Source:

- NYC covid-19:  
<https://data.cityofnewyork.us/Health/COVID-19-Daily-Counts-of-Cases-Hospitalizations-an/rc75-m7u3>
- Energy and water:  
<https://data.cityofnewyork.us/Environment/Energy-and-Water-Data-Disclosure-for-Local-Law-84-/qb3v-bbre>
- Vehicle Collisions:  
<https://data.cityofnewyork.us/Public-Safety/Motor-Vehicle-Collisions-Crashes/h9gi-nx95>
- NYC death cases(by year, used to compare 2020):  
<https://data.cityofnewyork.us/Health/New-York-City-Leading-Causes-of-Death/jb7j-dtam>
- NYC prisoner:  
<https://data.cityofnewyork.us/Public-Safety/Daily-Inmates-In-Custody/7479-ugqb>
- Violence:  
<https://data.cityofnewyork.us/Public-Safety/NYPD-Complaint-Data-Current-Year-To-Date-/5uac-w243>
- Shooting:  
<https://data.cityofnewyork.us/Public-Safety/NYPD-Shooting-Incident-Data-Year-To-Date-/5ucz-vwe8>

# Data Transformation and Cleaning

It varies by the dataset, general steps are below:

1. Read csv files into spark data frames
2. Attribute selection
3. Convert date to date type with extraction of year, month, day.
4. transform data to what will be used on tables if needed.
5. sort by date.

## Process Covid19 data

Before transformation:

- ▶ (2) Spark Jobs
- ▶  rdd1: pyspark.sql.dataframe.DataFrame = [DATE\_OF\_INTEREST: string, CASE\_COUNT: integer ... 37 more fields]

```
root
|-- DATE_OF_INTEREST: string (nullable = true)
|-- CASE_COUNT: integer (nullable = true)
|-- HOSPITALIZED_COUNT: integer (nullable = true)
|-- DEATH_COUNT: integer (nullable = true)
|-- DEATH_COUNT_PROBABLE      : integer (nullable = true)
|-- CASE_COUNT_7DAY_AVG      : integer (nullable = true)
|-- HOSP_COUNT_7DAY_AVG: integer (nullable = true)
|-- DEATH_COUNT_7DAY_AVG: integer (nullable = true)
|-- BX_CASE_COUNT: integer (nullable = true)
|-- BX_HOSPITALIZED_COUNT: integer (nullable = true)
|-- BX_DEATH_COUNT: integer (nullable = true)
|-- BX_CASE_COUNT_7DAY_AVG: integer (nullable = true)
|-- BX_HOSPITALIZED_COUNT_7DAY_AVG: integer (nullable = true)
|-- BX_DEATH_COUNT_7DAY_AVG: integer (nullable = true)
|-- BK_CASE_COUNT: integer (nullable = true)
|-- BK_HOSPITALIZED_COUNT: integer (nullable = true)
|-- BK_DEATH_COUNT: integer (nullable = true)
|-- BK_CASE_COUNT_7DAY_AVG: integer (nullable = true)
|-- BK_HOSPITALIZED_COUNT_7DAY_AVG: integer (nullable = true)
|-- BK_DEATH_COUNT_7DAY_AVG: integer (nullable = true)
```

Command took 0.91 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:40:41 PM on cs779 project

After transformation:

```
1 # only keep date, cases, hospitalized and death and sort by date
2 col_covid = ['DATE', 'CASE_COUNT', 'HOSPITALIZED_COUNT', 'DEATH_COUNT']
3 covid = rdd1.withColumn('DATE', to_date('DATE_OF_INTEREST', 'MM/dd/yyyy')) .select(col_covid).sort('DATE', ascending=False)
4 print(covid.count())
5 covid.show()
6 covid.write.saveAsTable('COVID19')
```

▶ (6) Spark Jobs

▶ covid: pyspark.sql.dataframe.DataFrame = [DATE: date, CASE\_COUNT: integer ... 2 more fields]

266

DATE	CASE_COUNT	HOSPITALIZED_COUNT	DEATH_COUNT
2020-11-20	831	47	4
2020-11-19	1169	88	3
2020-11-18	1282	95	7
2020-11-17	1337	78	9
2020-11-16	1589	110	9
2020-11-15	816	82	12
2020-11-14	947	81	10
2020-11-13	1418	75	6
2020-11-12	1408	65	14
2020-11-11	1436	78	9
2020-11-10	1508	66	7
2020-11-09	1511	73	4
2020-11-08	761	51	9
2020-11-07	798	55	14
2020-11-06	1002	53	10
2020-11-05	1104	71	10
2020-11-04	1077	54	9

Command took 4.34 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:48:03 PM on cs779 project

1. Extracted needed attributes.
2. Sort data by date.
3. Save as a table "Covid19".

## Process Shooting Incidents data

Before:

```

1 # preprocess shooting data
2 rdd2 = spark.read.csv('/FileStore/tables/NYPD_Shooting_Incident_Data__Year_To_Date_.csv', header = True, inferSchema=True)
3 rdd2.printSchema()

```

▶ (2) Spark Jobs

▶  rdd2: pyspark.sql.dataframe.DataFrame = [INCIDENT\_KEY: integer, OCCUR\_DATE: string ... 17 more fields]

```

root
|-- INCIDENT_KEY: integer (nullable = true)
|-- OCCUR_DATE: string (nullable = true)
|-- OCCUR_TIME: string (nullable = true)
|-- BORO: string (nullable = true)
|-- PRECINCT: integer (nullable = true)
|-- JURISDICTION_CODE: integer (nullable = true)
|-- LOCATION_DESC: string (nullable = true)
|-- STATISTICAL_MURDER_FLAG: boolean (nullable = true)
|-- PERP_AGE_GROUP: string (nullable = true)
|-- PERP_SEX: string (nullable = true)
|-- PERP_RACE: string (nullable = true)
|-- VIC_AGE_GROUP: string (nullable = true)
|-- VIC_SEX: string (nullable = true)
|-- VIC_RACE: string (nullable = true)
|-- X_COORD_CD: integer (nullable = true)
|-- Y_COORD_CD: integer (nullable = true)
|-- Latitude: double (nullable = true)
|-- Longitude: double (nullable = true)
|-- New Georeferenced Column: string (nullable = true)

```

Command took 1.15 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:36:39 PM on cs779 project

After:

```

1 shooting = rdd2.withColumn('DATE', to_date('OCCUR_DATE', 'MM/dd/yyyy')).select(['DATE', 'BORO', 'LOCATION_DESC', 'PERP_AGE_GROUP', 'PERP_SEX',
2 'STATISTICAL_MURDER_FLAG']).sort('DATE', ascending = False)
3 shooting.printSchema()
4 shooting.show()
5 shooting.write.saveAsTable('SHOOTING')

```

▶ (4) Spark Jobs

▶  shooting: pyspark.sql.dataframe.DataFrame = [DATE: date, BORO: string ... 8 more fields]

```

root
|-- DATE: date (nullable = true)
|-- BORO: string (nullable = true)
|-- LOCATION_DESC: string (nullable = true)
|-- PERP_AGE_GROUP: string (nullable = true)
|-- PERP_SEX: string (nullable = true)
|-- PERP_RACE: string (nullable = true)
|-- VIC_AGE_GROUP: string (nullable = true)
|-- VIC_SEX: string (nullable = true)
|-- VIC_RACE: string (nullable = true)
|-- STATISTICAL_MURDER_FLAG: boolean (nullable = true)

```

DATE	BORO	LOCATION_DESC	PERP_AGE_GROUP	PERP_SEX	PERP_RACE	VIC_AGE_GROUP	VIC_SEX	VIC_RACE	STATISTICAL_MURDER_FLAG
2020-09-30	BROOKLYN	MULTI DWELL - APT...	18-24	M	BLACK	18-24	M	WHITE HISPANIC	false
2020-09-30	QUEENS	MULTI DWELL - APT...	25-44	M	WHITE HISPANIC	25-44	F	WHITE HISPANIC	true
2020-09-30	BROOKLYN	MULTI DWELL - PUB...	null	null	null	25-44	M	BLACK	false
2020-09-30	BROOKLYN	COMMERCIAL BLDG	null	null	null	18-24	M	UNKNOWN	false
2020-09-30	BROOKLYN	COMMERCIAL BLDG	null	null	null	18-24	M	UNKNOWN	true
2020-09-30	BROOKLYN	COMMERCIAL BLDG	null	null	null	25-44	M	UNKNOWN	true

Command took 4.45 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:40:12 PM on cs779 project

1. Attribute selection.
2. Convert date type and sort by date.
3. Save as a table "Shooting".

# Process Prisoner data

Before:

1 %python

2 # preprocess prisoner data

3

4 rdd3 = spark.read.csv('/FileStore/tables/Daily\_Inmates\_In\_Custody.csv', header = True, inferSchema=True)

5 rdd3.printSchema()

6 rdd3.show()

▶ (3) Spark Jobs

▶ rdd3: pyspark.sql.dataframe.DataFrame = [INMATEID: integer, ADMITTED\_DT: string ... 11 more fields]

INMATEID	ADMITTED_DT	DISCHARGED_DT	CUSTODY_LEVEL	BRADH	RACE	GENDER	AGE	INMATE_STATUS_CODE	SEALED	SRG_FLG	TOP_CHARGE	INFRACTION
20149472	09/14/2020 05:02:...	null	MIN	Y	B	M	30	DE	N	N	120.10	N
33284	11/20/2020 02:52:...	null	MED	N	B	M	39	DPV	N	N	null	N
20029823	01/06/2019 12:24:...	null	MAX	Y	B	M	36	DE	N	N	135.20	Y
20211644	10/15/2020 01:42:...	null	MED	N	W	M	52	DE	N	N	220.21	N
20202151	11/17/2020 12:41:...	null	MED	N	O	M	20	DE	N	N	265.03	N
20019267	03/14/2020 12:29:...	null	MAX	Y	B	M	28	DE	N	Y	120.10	Y
20009388	10/13/2020 12:47:...	null	MAX	Y	B	M	32	DEP	N	N	160.10	N
37060	11/03/2020 10:40:...	null	MED	Y	B	M	31	DPV	N	N	null	N
20036616	06/24/2019 01:23:...	null	MAX	Y	B	M	39	DE	N	Y	230.34	Y
20196814	07/19/2018 11:54:...	null	MAX	N	B	M	21	DE	N	Y	125.25	Y
20126766	12/06/2019 05:32:...	null	MED	N	B	M	29	DE	N	N	130.96	Y
20211995	11/12/2020 12:19:...	null	MIN	N	W	M	18	DE	N	N	155.35	N
20192843	09/27/2020 12:46:...	null	MAX	N	W	M	36	DE	N	N	160.15	N
20207687	01/25/2020 12:21:...	null	MAX	Y	W	M	28	DE	N	Y	125.25	Y
20211077	09/06/2020 12:58:...	null	MIN	Y	O	M	28	DE	N	N	150.20	N
20209165	02/01/2020 01:15:...	null	MED	Y	A	M	25	DE	N	N	125.25	N
20017335	08/05/2019 12:18:...	null	MED	Y	B	M	28	DE	N	N	140.30	N
20003013	10/08/2020 05:23:...	null	MAX	Y	B	F	31	DE	N	N	110-160.05	Y
20200525	12/09/2018 10:08:...	null	MED	Y	B	M	60	DE	N	N	125.25	N

Command took 1.95 seconds -- by frankwc6@bu.edu at 11/25/2020, 3:05:06 PM on CS799\_Project

After:

```

1 %python
2 # extract date only without time of the day
3 string_udf = F.udf(lambda x: x[:10])
4 prisoner = rdd3.withColumn('DATE', F.to_date(string_udf('ADMITTED_DT'), 'MM/dd/yyyy')).select(['DATE', 'CUSTODY_LEVEL', 'BRADH', 'GENC
5 prisoner.show()
6 prisoner.write.saveAsTable('PRISONER')

```

▶ (4) Spark Jobs

▶ prisoner: pyspark.sql.dataframe.DataFrame = [DATE: date, CUSTODY\_LEVEL: string ... 5 more fields]

DATE	CUSTODY_LEVEL	BRADH	GENDER	AGE	RACE	INFRACTION
2020-11-23	null	N	M	58	O	N
2020-11-23	null	N	M	33	B	N
2020-11-23	null	N	M	30	B	N
2020-11-23	null	N	M	32	B	N
2020-11-23	null	N	M	37	B	N
2020-11-23	null	N	M	43	B	N
2020-11-23	null	N	M	36	B	N
2020-11-23	null	N	M	43	B	N
2020-11-22	null	N	M	31	W	N
2020-11-22	null	N	M	37	B	N
2020-11-22	null	N	M	26	B	N
2020-11-22	null	N	M	39	B	N
2020-11-22	null	N	M	26	O	N
2020-11-22	null	N	M	26	B	N
2020-11-22	null	N	M	21	B	N
2020-11-22	MAX	N	M	28	B	N
2020-11-22	null	N	M	27	B	N
2020-11-22	null	N	M	31	B	N

Command took 5.68 seconds -- by frankwc6@bu.edu at 11/25/2020, 4:44:59 PM on CS799\_Project

1. Extract date as “yyyy-mm-dd” format without day time.
2. Attribute selection.
3. Sort by date and save as table “Prisoner”

## Process utility data

Data and metrics on water and energy consumption in buildings over 25,000 ft2. The original dataset has more than 60 attributes, I use only averaged gas, water, electricity consumption, emission and building usage. Numeric attributes were divided by the number of occupants to get the means of those columns.



## Code and the final table:

```
1 # process utility data
2 rdd4 = spark.read.csv('/FileStore/tables/Energy_and_Water_Data_Disclosure_for_Local_Law_84_2020__Data_for_Calendar_Year_2019_.csv', header=True, inferSchema=True )
3 # check column names
4 cols = rdd4.columns
5 # pick up needed columns from more than 60 columns
6 index = [60, 14, 17, 27, 48, 51, 54, 57]
7 columns = [cols[i] for i in index]
8 # select those columns and check if they are correct
9 rdd4 = rdd4.select(columns)
10 rdd4.columns
```

```
↳ (2) Spark Jobs
↳ rdd4: pyspark.sql.dataframe.DataFrame = [Generation Date: string, Borough: string ... 6 more fields]
```

```
Out[1]: ['Generation Date',
' Borough',
' Primary Property Type - Self Selected',
' Occupancy',
' Weather Normalized Site Natural Gas Use (therms)',
' Weather Normalized Site Electricity (kWh)',
' Total GHG Emissions (Metric Tons CO2e)',
' Water Use (All Water Sources) (kgal)']
```

Command took 19.57 seconds -- by: frankac6@bu.edu at 12/4/2020, 2:29:05 PM on cs779 project

Cell 8

```
1 # change column names to be short and convert date
2 df = rdd4.select(col('Generation Date').alias('Date'), col('Primary Property Type - Self Selected').alias('Usage'), col('Occupancy'), col('Weather Normalized Site Natural Gas Use (therms)').alias('Gas'),
3 col('Weather Normalized Site Electricity (kWh)').alias('Electricity'), col('Total GHG Emissions (Metric Tons CO2e)').alias('Emission'), col('Water Use (All Water Sources) (kgal)').alias('Water'))
4 df = df.withColumn('DATE', to_date('Date', 'MM/dd/yyyy')) .sort('DATE', ascending = False)
5 df.printSchema()
6 df.show()
```

```
1 # clean data and format it as one person usage
2 # an issue is that multiple null value representations, it causes dataframe could do computation and hard to convert types, so I use RDD instead.
3 def isfloat(x):
4     """return float values"""
5     try:
6         float(x)
7         return True
8     except:
9         return False
10
11
12 def filter_value(x):
13     """filter float values"""
14     if isfloat(x[2]) and x[2] != 0 and isfloat(x[3]) and isfloat(x[4]) and isfloat(x[5]) and isfloat(x[6]):
15         return x
16
17
18 def get_ave(x):
19     """divide gas, electricity, emission and water by the number of occupants"""
20     return (x[0], x[1], x[2], round(float(x[3]) / float(x[2]), 2), round(float(x[4]) / float(x[2]), 2), round(float(x[5]) / x[2], 2), round(float(x[6]) / x[2], 2) )
21
22
23 rdd_df = df.rdd
24 print('Before filtering:', rdd_df.count())
25 header = rdd_df.first()
26 rdd_df = rdd_df.subtract(sc.parallelize([header])).map(tuple)
27 rdd_df = rdd_df.filter(filter_value)
28 print('After filtering:', rdd_df.count())
29 rdd_df = rdd_df.map(lambda x: get_ave(x) )
30 rdd_df.take(10)
31 utility_df = spark.createDataFrame(rdd_df, ['DATE', 'Usage', 'Occupancy', 'Gas', 'Electricity', 'Emission', 'Water'])
32 #utility_df.show()
33 utility_df.printSchema()
34 utility_df.show()
35 utility_df.coalesce(1).write.saveAsTable('UTILITY')
```

```
↳ (6) Spark Jobs
↳ utility_df: pyspark.sql.dataframe.DataFrame = [DATE: date, Usage: string ... 5 more fields]
```

DATE	Usage	Occupancy	Gas	Electricity	Emission	Water
[2020-11-03]	[Multifamily Housing]	[95]	[254.76]	[3986.83]	[2.45]	[39.15]
[2020-11-02]	[Hotel]	[108]	[398.39]	[10395.76]	[4.93]	[41.46]
[2020-11-02]	[Multifamily Housing]	[95]	[555.19]	[1927.4]	[3.49]	[28.04]
[2020-10-31]	[Multifamily Housing]	[98]	[56.21]	[1649.6]	[0.89]	[0.95]

## Process complaint data

This dataset includes all valid felony, misdemeanor, and violation crimes reported to the New York City Police Department (NYPD).

Before:

```

1 # process complaints data which includes conflicts and crimes
2
3 rdd5 = spark.read.csv('/FileStore/tables/NYPD_Complaint_Data_Current__Year_To_Date.csv', header=True, inferSchema=True )
4 rdd5.printSchema()
5 rdd5.show()

```

► (3) Spark Jobs

► rdd5: pyspark.sql.dataframe.DataFrame = [CMPLNT\_NUM: integer, ADDR\_PCT\_CD: integer ... 34 more fields]

CMPLNT_NUM	ADDR_PCT_CD	BORO_NM	CMPLNT_FR_DT	CMPLNT_FR_TM	CMPLNT_TO_DT	CMPLNT_TO_TM	CRM_ATPT_CPTD_CD	HADEVELOPT	HOUSING_PS
OFNS_DESC	PARKS_NM	PATROL_BORO	PD_CD	PD_DESC	PREM_TYP_DESC	RPT_DT	STATION_NAME	SUSP_AGE_GROUP	SUSP_RACE
ORD_CD	Latitude	Longitude	Lat_Lon	New Georeferenced Column					
972326799	81	null	09/28/2020	21:27:00	null	null	COMPLETED	null	nul
DER & NON-NEGL...	null	null	null	null	null	null	09/28/2020	null	null
904	186483	40.67851591200008	-73.92914304899993	(40.6785159120000...	POINT (-73.929143...				
376304873	52	null	09/27/2020	19:13:00	null	null	COMPLETED	null	nul
DER & NON-NEGL...	null	null	null	null	null	null	09/27/2020	null	45-64 WHITE H
928	258050	40.87490600500007	-73.87822380899996	(40.8749060050000...	POINT (-73.878223...				
299326203	75	null	09/21/2020	01:21:00	null	null	COMPLETED	null	nul
DER & NON-NEGL...	null	null	null	null	null	null	09/21/2020	null	null
234	181211	40.66399002800006	-73.86669235099998	(40.6639900280000...	POINT (-73.866692...				
674946147	121	null	09/15/2020	08:46:00	null	null	COMPLETED	null	nul
DER & NON-NEGL...	null	null	null	null	null	null	09/15/2020	null	null
430	170972	40.63584491100005	-74.165090337	(40.6358449110000...	POINT (-74.165090...				
416422620	101	null	09/08/2020	13:50:00	null	null	COMPLETED	null	nul
DER & NON-NEGL...	null	null	null	null	null	null	09/08/2020	null	<18
837	157548	40.59887464700005	-73.76382298499993	(40.5988746470000...	POINT (-73.763822...				

Command took 5.89 seconds -- by frankwc6@bu.edu at 11/26/2020, 2:50:56 PM on cs779 project

After:

```

1 cols = ['DATE', 'OFNS_DESC', 'SUSP_AGE_GROUP', 'SUSP_RACE', 'SUSP_SEX', 'VIC_AGE_GROUP', 'VIC_RACE', 'VIC_SEX']
2 rdd5 = rdd5.withColumn('DATE', to_date('CMPLNT_FR_DT', 'MM/dd/yyyy')) .select(cols).sort('DATE', ascending = False)
3 rdd5.show()
4 # rdd5.write.saveAsTable('CRIME')
5 rdd5.printSchema()
6 # filter incomplete rows
7 file_read = rdd5.filter((rdd5['SUSP_AGE_GROUP'] != 'UNKNOWN') & (rdd5['SUSP_RACE'] != 'UNKNOWN') &
8                       (rdd5['SUSP_SEX'] != 'UNKNOWN') & (rdd5['VIC_AGE_GROUP'] != 'UNKNOWN') & (rdd5['VIC_RACE'] != 'UNKNOWN') &
9                       (rdd5['VIC_SEX'] != 'UNKNOWN'))
10 # a problem is the data too big to be converted to a table, so I saved it to a csv file and then read the file into a table
11 file_read.write.csv('crime.csv', header = True)

```

► (4) Spark Jobs

► rdd5: pyspark.sql.dataframe.DataFrame = [DATE: date, OFNS\_DESC: string ... 6 more fields]

► file\_read: pyspark.sql.dataframe.DataFrame = [DATE: date, OFNS\_DESC: string ... 6 more fields]

DATE	OFNS_DESC	SUSP_AGE_GROUP	SUSP_RACE	SUSP_SEX	VIC_AGE_GROUP	VIC_RACE	VIC_SEX
2020-09-30	CRIMINAL MISCHIEF...	UNKNOWN	UNKNOWN	U	25-44	WHITE HISPANIC	F
2020-09-30	FELONY ASSAULT	25-44	BLACK	F	UNKNOWN	UNKNOWN	E
2020-09-30	ASSAULT 3 & RELAT...	25-44	BLACK	M	25-44	BLACK	F
2020-09-30	PETIT LARCENY	null	null	null	UNKNOWN	UNKNOWN	D
2020-09-30	BURGLARY	UNKNOWN BLACK	HISPANIC	M	UNKNOWN BLACK	HISPANIC	D
2020-09-30	ASSAULT 3 & RELAT...	UNKNOWN	BLACK	F	45-64	WHITE HISPANIC	F
2020-09-30	GRAND LARCENY	UNKNOWN	UNKNOWN	M	25-44	WHITE	F
2020-09-30	ASSAULT 3 & RELAT...	UNKNOWN	BLACK	M	45-64	BLACK	M
2020-09-30	PETIT LARCENY	UNKNOWN	UNKNOWN	M	UNKNOWN	UNKNOWN	D
2020-09-30	PETIT LARCENY	null	null	null	UNKNOWN	UNKNOWN	D
2020-09-30	BURGLARY	UNKNOWN	BLACK	M	UNKNOWN	UNKNOWN	D
2020-09-30	GRAND LARCENY OF ...	null	null	null	UNKNOWN	UNKNOWN	D
2020-09-30	PROSTITUTION & RE...	<18 BLACK	HISPANIC	M	UNKNOWN	UNKNOWN	E
2020-09-30	NYS LAWS-UNCLASSI...	25-44	BLACK	M	UNKNOWN	UNKNOWN	E
2020-09-30	DANGEROUS WEAPONS	null	null	null	UNKNOWN	UNKNOWN	E
2020-09-30	PETIT LARCENY	UNKNOWN	UNKNOWN	U	25-44	BLACK	F
2020-09-30	FELONY ASSAULT	25-44	BLACK	M	25-44	BLACK	F
2020-09-30	ROBBERY	UNKNOWN	BLACK	M	UNKNOWN	UNKNOWN	M

Command took 16.89 seconds -- by frankwc6@bu.edu at 11/26/2020, 2:51:03 PM on cs779 project

## Process Vehicle Collision data

The Motor Vehicle Collisions crash table contains details on the crash event. Each row represents a crash event. The Motor Vehicle Collisions data tables contain information from all police reported motor vehicle collisions in NYC.

Before:

```
1 rdd6 = spark.read.csv('/FileStore/tables/Motor_Vehicle_Collisions___Crashes.csv', inferSchema=True, header= True)
2 rdd6.printSchema()
3 rdd6.show()
```

▶ (3) Spark Jobs

▶ rdd6: pyspark.sql.dataframe.DataFrame = [CRASH DATE: string, CRASH TIME: string ... 27 more fields]

CRASH DATE	CRASH TIME	BOROUGH	ZIP CODE	LATITUDE	LONGITUDE	LOCATION	ON STREET NAME	CROSS STREET NAME	NUMBER OF PERSONS INJURED	NUMBER OF PERSONS KILLED	NUMBER OF CYCLIST INJURED	NUMBER OF CYCLIST KILLED	NUMBER OF MOTORIST INJURED	NUMBER OF MOTORIST KILLED	CONTRIBUTING FACTOR VEHICLE 1	CONTRIBUTING FACTOR VEHICLE 2	CONTRIBUTING FACTOR VEHICLE 3	CONTRIBUTING FACTOR VEHICLE 4	CONTRIBUTING FACTOR VEHICLE 5	COLLISION_ID	VEHICLE TYPE CODE 1	VEHICLE TYPE CODE 2
06/22/2020	14:00	MANHATTAN	10019	40.770523	-73.99196	(40.770523, -73.99196)			0	0	0	0	0	0								606
08/02/2020	23:20	QUEENS	11385	40.701683	-73.90885	(40.701683, -73.90885)	Box Truck		0	0	0	0	0	0								
07/12/2020	18:45	MANHATTAN	10040	40.855133	-73.93688	(40.855133, -73.93688)	Sedan	Station Wagon/Spo...	0	0	0	0	0	0								
08/06/2020	19:16	QUEENS	11362	40.760284	-73.73177	(40.760284, -73.73177)			0	0	0	0	0	0								248-7
06/23/2020	15:41	QUEENS	11432	40.708378	-73.79169	(40.708378, -73.79169)	Station Wagon/Spo...		0	0	0	0	0	0								

Command took 17.70 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:04:52 PM on cs779 project

After:

```
1 incident = rdd6.groupBy('CRASH DATE', 'BOROUGH').agg(F.count('NUMBER OF PERSONS INJURED').alias('injured_count'),
2 F.count('NUMBER OF PERSONS KILLED').alias('death_count')).withColumn('Date',
3 to_date('CRASH DATE', 'MM/dd/yyyy')).select(['DATE', 'BOROUGH', 'injured_count', 'death_count']).sort('DATE', ascending = False)
4 incident.show()
5 incident.printSchema()
```

▶ (2) Spark Jobs

▶ incident: pyspark.sql.dataframe.DataFrame = [DATE: date, BOROUGH: string ... 2 more fields]

DATE	BOROUGH	injured_count	death_count
2020-11-20	BROOKLYN	70	70
2020-11-20	MANHATTAN	20	20
2020-11-20	STATEN ISLAND	3	3
2020-11-20	QUEENS	53	53
2020-11-20	BRONX	29	29
2020-11-20	null	87	87
2020-11-19	QUEENS	42	42
2020-11-19	MANHATTAN	12	12
2020-11-19	STATEN ISLAND	3	3
2020-11-19	BROOKLYN	52	52
2020-11-19	BRONX	32	32
2020-11-19	null	84	84
2020-11-18	MANHATTAN	23	23
2020-11-18	null	84	84
2020-11-18	BRONX	36	36
2020-11-18	QUEENS	44	44
2020-11-18	BROOKLYN	67	67
2020-11-18	STATEN ISLAND	3	3

Command took 13.41 seconds -- by frankwc6@bu.edu at 11/26/2020, 3:20:02 PM on cs779 project

I summarized this dataset by counting the number of injuries and deaths. A finding is injury and death are almost the same, it shows the fatality rate of vehicle collision is close to 100%

# Construct Tables

1. Save data frames into tables.
2. Save large data frames into csv files and read files by tables.
3. Summarize tables into views as dimension tables.
4. Load summarized data into a fact table. I'm not sure this can be called to be a fact table because those dimensions only share year, month and day.
5. Do queries. The query pattern is: fact table--- views --- original table. This could improve performance significantly if we want to find out a specific question.

The most complicated view was the utility view. A big problem was gas, water, electricity and emission are not on the same scale. So, I normalized those values by using the function of  $(\text{value} - \text{min}) / (\text{max} - \text{min})$ . This function transforms all values to around 1. After getting all 1, I scaled them up by multiplying 10000.

```
%sql
--- get average daily average utility usage

---CREATE VIEW bridge_table as

with temp_table(date, gas_ave, electricity_ave, water_ave, emission_ave) as (
select date, sum(Gas)/ count(*) as gas_ave, sum(Electricity) / count(*) as electricity_ave, sum(Water)/ count(*) as water_ave, sum(Emission) / count(*) as emission_ave
from utility
group by date order by date desc)

select min(gas_ave) as min_gas, max(gas_ave) as max_gas, min(electricity_ave) as min_electricity, max(electricity_ave) as max_electricity, min(water_ave) as min_water, max(water_ave) as max_water,
min(emission_ave) as min_emission, max(emission_ave) as max_emission from temp_table
```

(3) Spark Jobs

	min_gas	max_gas	min_electricity	max_electricity	min_water	max_water	min_emission	max_emission
1	16.59	14301.895384615385	388.44	185741.84000000003	3.23	1030704.230882353	1.19	131.47999999999996

Showing all 1 rows.

mand took 10.78 seconds -- by frankwc@bu.edu at 12/5/2020, 2:30:28 PM on CS779

```
Cmd 24
1 %sql
2
3 --- get normalize those data and scale up to make them comparable
4 --CREATE VIEW NORMALIZED_UTILITY AS
5
6 with temp_table(date, gas_ave, electricity_ave, water_ave, emission_ave) as (
7 select date, round(sum(gas)/ count(*), 2) as gas_ave, round(sum(electricity) / count(*), 2) as electricity_ave, round(sum(water)/ count(*), 2) as water_ave, sum(emission) / count(*) as emission_ave
8 from utility
9 group by date order by date desc)
10
11
12 select year(date) as year, month(date) as month, day(date) as day, round((gas_ave - min_gas) / (max_gas - min_gas) * 10000) as norm_gas, round((electricity_ave - min_electricity) / (max_electricity - min_electricity) * 10000) as norm_electricity, round((water_ave - min_water) / (max_water - min_water) * 10000) as norm_water, round((emission_ave - min_emission) / (max_emission - min_emission) * 10000) as norm_emission from temp_table, bridge_table order by year desc, month desc, day desc

OK
Command took 1.41 seconds -- by frankc@bu.edu at 11/30/2020, 7:25:27 PM on cs779 (clone)

Cmd 25
1 %sql
2 --- check normalized utility view
3 select month, sum(norm_gas) as gas, sum(norm_electricity) as electricity, sum(norm_water) as water, sum(norm_emission) as emission from NORMALIZED_UTILITY group by month order by month desc

(8) Spark Jobs
```

	month	gas	electricity	water	emission
1	11	2070	4722	3	2644
2	10	14437	17930	43	14823
3	9	25863	22090	10	21403
4	8	17452	26465	35	20620
5	7	11332	15380	19	12728
6	6	10607	13864	14	11562
7	5	11353	19906	25	24608

Showing all 9 rows.

For other views were only transformation to daily counts of certain attributes.

The last step is to join all tables to a new summary table. The final table only has data from March 2020 to September 2020. As the primary of this project, it focuses on data from the beginning of covid to the end of intersection among all views. After getting the summary table, we can do queries now.

```
Cmd 30
1 %sql
2
3 --- the table I will use with intersection date of all tables
4 --create table fact_table as
5
6 select P.year, P.month, P.day, prisoner_sum, shooting_case_count, CASE_COUNT, HOSPITALIZED_COUNT, DEATH_COUNT, cases as crime_daily,
7 norm_gas, norm_electricity, norm_water, norm_emission, daily_injured as vehicle_collision_injured, daily_death as vehicle_collision_death
8 from prisoner_by_day P
9 join covid19_view C on P.year = C.year and P.month = C.month and P.day = C.day
10 join daily_shooting_view D on P.year = D.year and P.month = D.month and P.day = D.day
11 join crime_view M on P.year = M.year and P.month = M.month and P.day = M.day
12 join normalized_utility N on P.year = N.year and P.month = N.month and P.day = N.day
13 join vehicle_collisions_view V on P.year = V.year and P.month = V.month and P.day = V.day
14 order by year desc, month desc, day desc

Cmd 31
1 %sql
2
3 select * from fact_table

(1) Spark Jobs
```

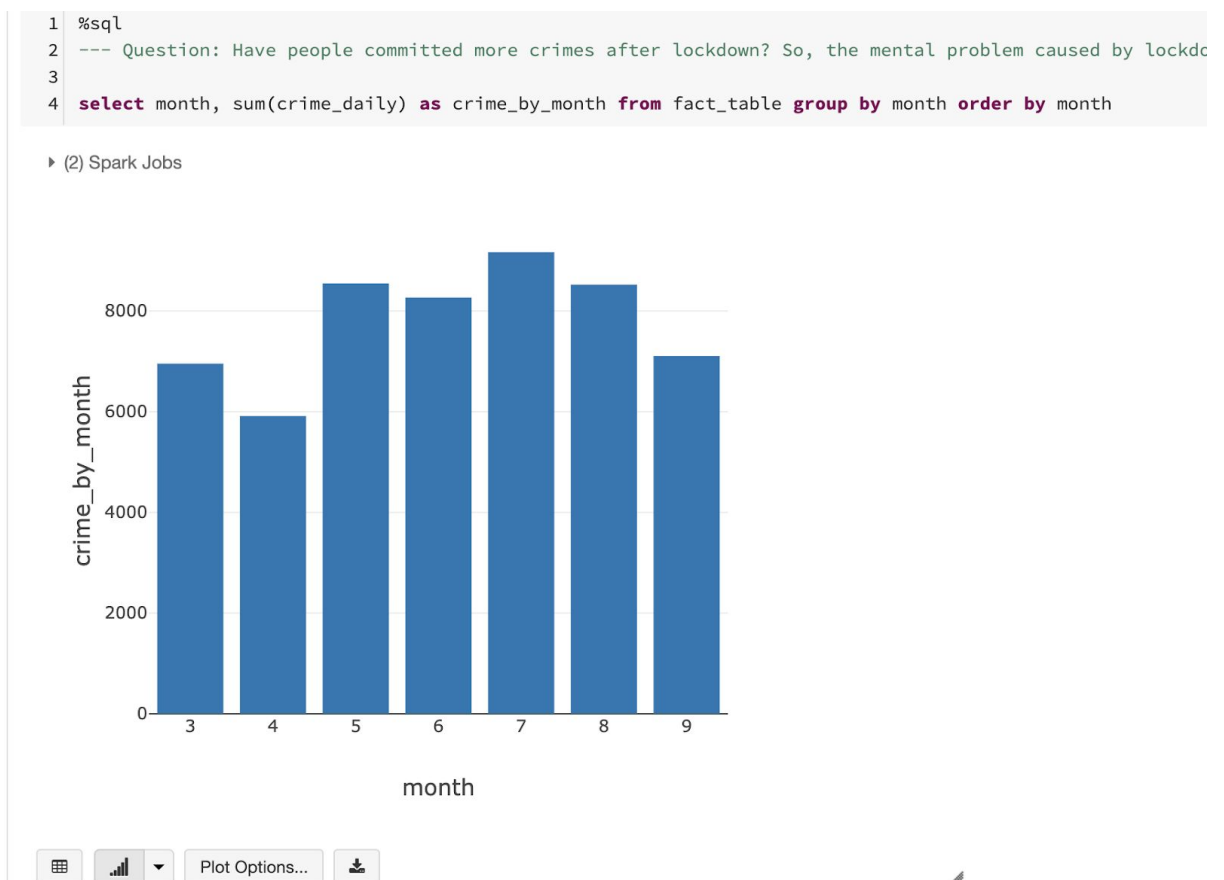
	year	month	day	prisoner_sum	shooting_case_count	CASE_COUNT	HOSPITALIZED_COUNT	DEATH_COUNT	crime_daily	norm_gas	norm_electricity	norm_water	norm_emission
1	2020	9	30	26	8	588	51	7	225	203	166	0	139
2	2020	9	29	23	2	688	49	8	288	737	777	0	690
3	2020	9	28	15	4	435	36	4	305	1379	185	0	832
4	2020	9	25	18	5	454	49	4	313	153	235	0	150
5	2020	9	24	24	4	468	33	3	289	50	178	0	191
6	2020	9	23	29	7	533	43	4	273	347	565	1	360
7	2020	9	22	12	2	407	32	3	287	168	181	0	91

Showing all 173 rows.

## Questions to explore

1. Have people committed more crimes after lockdown? So, the mental problem caused by lockdown made public safety worse?


We first see the crime rate by month in 2020 from the summary table. It shows a decline from March to April and then bounced back in June. This change shows the decline of crime during lockdown but did not continue after reopening the country. Now, we need to know if covid19 changed the crime by a social impact but not the lockdown time.



Let's check if any category of crime has been changed by covid19.

```
1 %sql
2
3 select OFNS_DESC, month(date) as month, count(*) as count from crime where year(date) = 2020 group by OFNS_DESC, month order by month
4
5 --- Answer: No, they are the same before and after covid lockdown
```

▶ (2) Spark Jobs



Only showing the first twenty series.

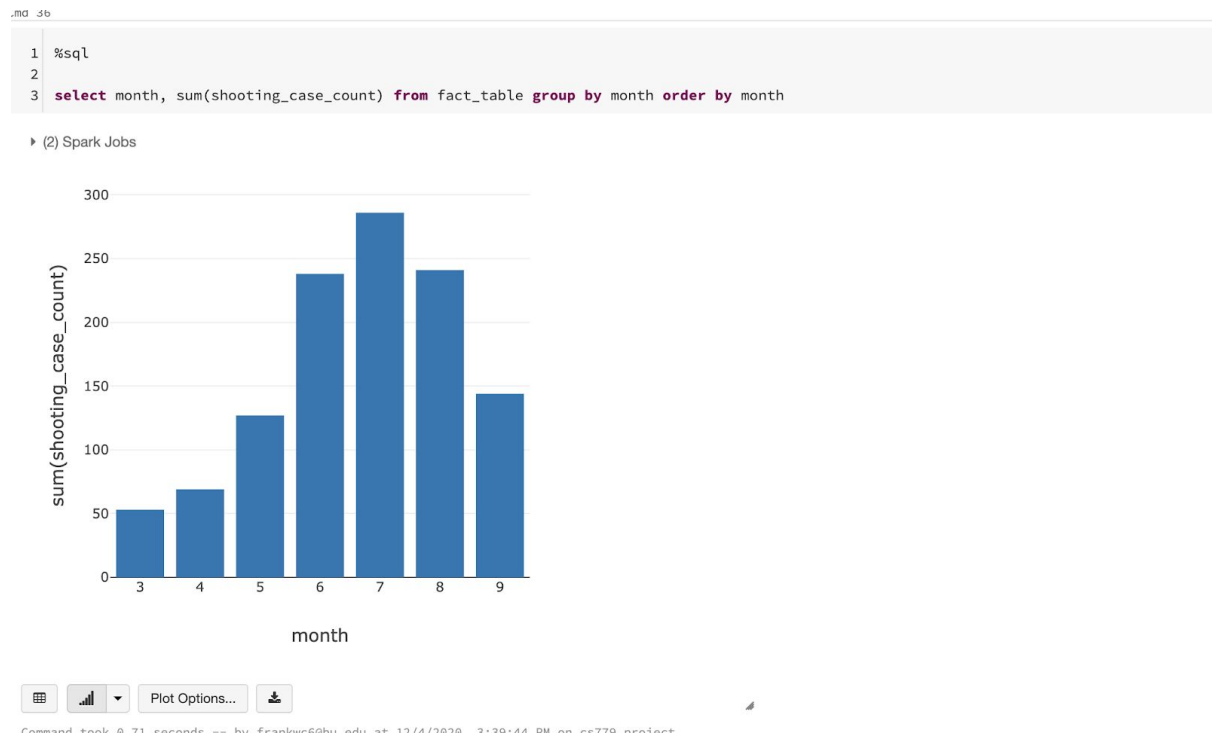
The plot shows almost the same count before and after lockdown. So, we can conclude that Covid19 did not change the crime rate in New York city.

## 2. If the crime rate didn't change, what about shooting cases?

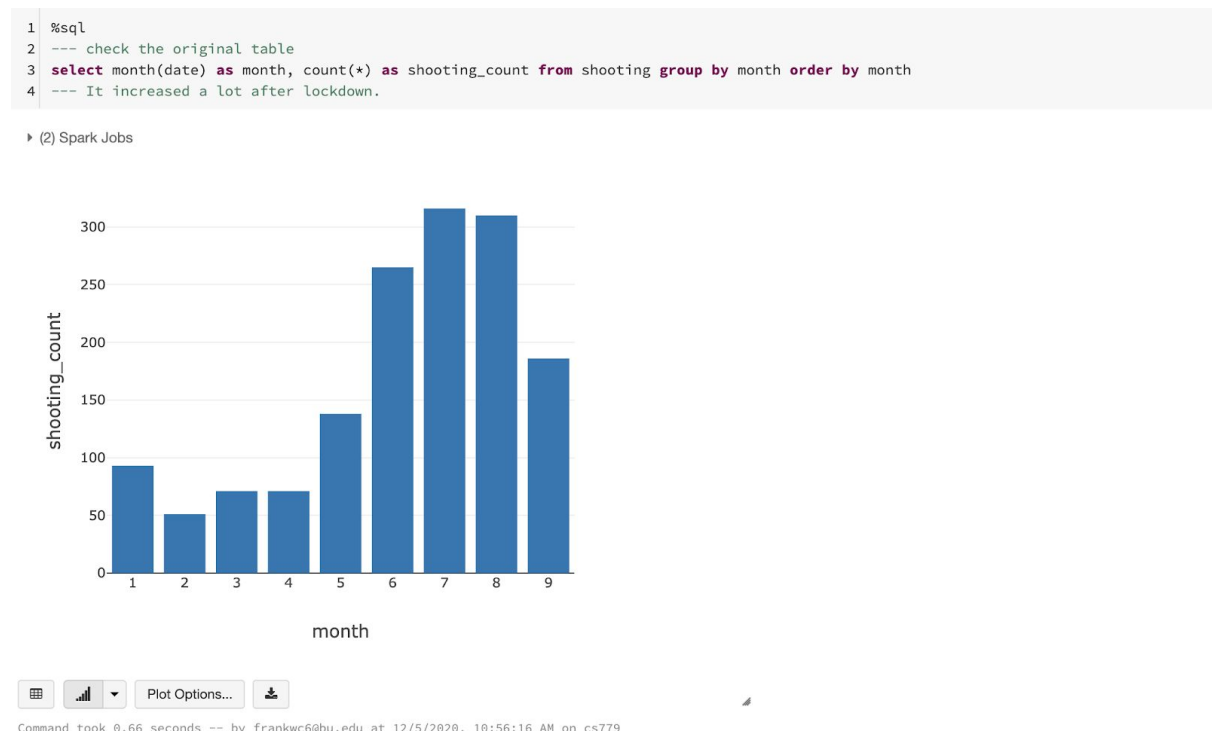
A thing should be noticed is the crime dataset exclude shooting cases, they are general cases. So, I will keep digging into the question of have covid19 made more shooting cases.

The bar plot shows a significant increase in shooting cases after April 2020.





To confirm the result of increasing shooting cases after lockdown, we check the original shooting cases table, which has data from January to September. It showed the same trend of shooting cases increasing after lockdown.





But January was affected by covid19. Since we don't have historic data of shooting cases, we can't conclude that covid19 made shooting cases more by causing mental problems. There might be a seasonal reason or data integrity issue.

### 3. Has covid19 reduced energy consumption?




We can see a drop in March and April and then a recovery from May, even more consumption after lockdown. For this result, air conditionings were widely used in the summer, which are energy-consuming. Especially we use data from buildings with big size, they are mostly office buildings or shopping malls. Let's check the whole year.

```

1 %sql
2
3 select month(date) as month, round(sum(gas), 2) gas, round(sum(Electricity), 2) electricity, round(sum(Water), 2) water, round(sum(Emis
4
5 --- covid does reduced energy consumption, business bounced back in july but some of them deaded after that. Or the reason is people wo

```

▶ (2) Spark Jobs



The assumption is true, energy usage dropped after summer. We can't conclude the drop of energy usage was caused by covid19 because people have been working at home since March. I need to get more data for the past years.

## Challenges

1. Figuring out advantages and disadvantages of spark dataframe and SQL both, the data transformation process was based on the structure of tables. The tables serve for query use.
2. Transforming data is not as easy as cleaning data. I had to dissect query plans to the data needed. To make data comparable, I normalized the utility table.
3. Query plans could be hieratical. For example, I want to know whether the total crime cases increased and which one increased during covid-19. I should go to the fact table to get

overall statistics by month and go to the summary crime view to find categories. For some questions, we need to go back to the original tables to find relative attributes, such as streets, location descriptions and more information about criminals.

## Conclusion

1. Spark is designed like a distributed database, data was stored on multiple nodes with replicas. It has more flexibility for data manipulation by using dataframes but less efficient for doing queries. It does not as fast as querying a table.
2. A relational database could keep data consistent. This made queries much easier and prevented data from getting corrupted and easy to connect data.
3. To organize data and extract intelligence is beyond simple processing. This is one of the reasons why designing a data warehouse is difficult. Because it should be designed to serve analysis and it will be hard to make changes once massive data is loaded into a data warehouse.